

How to Resolve Blocking and Deadlock in SQL Server

Ying Fan

Support Engineer
Customer Service & Support
Asia Pacific & Greater China Region

Welcome!

Last Workshop Re-cap

- Index structure
- Reading Query Plans
- Troubleshooting Query Plans
- Optimizing Query Plans
- Plan Caching issue

Agenda

- Blocking Issue
 - Locking
 - PageIOlatch
 - Pagelatch
 - ASYNC NETWORK IO
 - Tempdb Contention
- Deadlock

BLOCKING

What is Blocking?

- An existing spid has a lock on the resource
- Your connection is requesting a lock on the same resource with an incompatible lock mode

Information collection

- Sys.sysprocesses
- Sys.dm_exec_requests
- Sys.dm_os_wait_stats
- Blocking script (KB271509)
- SQL trace
- Pssdiag

Information From Sysprocesses

- Blocked – what spid is blocking yours?
- Cmd – basic type of command being processes
- Status – is the spid doing anything
- Waittype – hex value denoting what the spid is waiting on
- Lastwaittype – saved each time a wait finishes
- Waittime – how long has it been waiting (milliseconds)
- Waitresource – text description of resource
- Open_tran – number of nested transaction

DMVs

- `sys.dm_exec_requests`
 - New DMV introduced from SQL Server 2005
 - `sys.sysprocesses` still available for backward compatibility purpose
 - Handle of the TSQL and query plan
- `sys.dm_exec_connections`
- `sys.dm_exec_sessions`

Waittype in Sys.dm_os_wait_stats

- LCX_M_x
- PAGEIOLATCH_x
- PAGELATCH_x
- ASYNC_NETWORK_IO
- WRITELOG
- CXPACKET
- SQLTRACE_x

LCK_M_x

- Used to Provide Desired Level of Logical Consistency of Data
- Type of locks acquired depend on:
 - Isolation level
 - Lock cost
 - lock escalation
 - Lock hints
- Locking issue could be caused:
 - Long running query/sp
 - Improper use of transaction

Lock Resources and Hierarchy

- Common Resources

- RID
- Key
- Page
- Table

- Lock Hierarchy

- Acquire intent locks before traversing to the level below

Resource	Format	Example
Table	DatabaseID:ObjectID	TAB: 5:1
Page	DatabaseID:File:Page	PAG: 5:1:104
Key	DatabaseID:ObjectID:IndexID (Hash value for index key)	KEY: 5:1977058079:1 (02014f0bec4e)
RID	DatabaseID:FileID:PageID:Slot	RID: 5:1:104:3

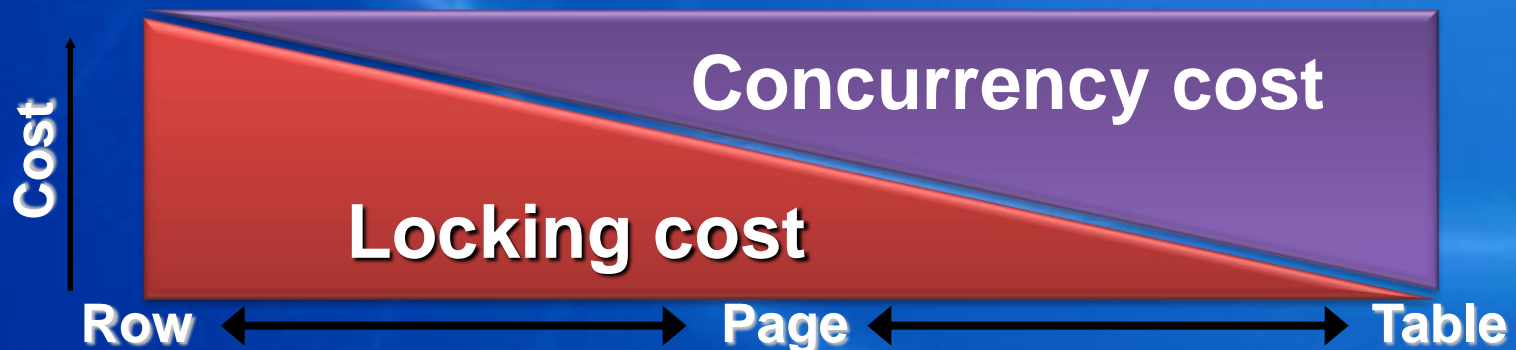
Lock Duration

- Lock Mode and Transaction Isolation Level

Lock Mode	Read Committed	Repeatable Read	Serializable
Shared	Held until data read and processed	Held until end of transaction	Held until end of transaction
Update	Held until data read and processed unless promoted to Exclusive	Held until data read and processed unless promoted to Exclusive	Held until end of transaction unless promoted to Exclusive
Exclusive	Held until end of transaction	Held until end of transaction	Held until end of transaction

Dynamic Locking

- Row locking is not always the right choice
 - Scanning a table with 100 million rows means 100 million calls to the lock manager
- Sometimes page or table locking is the optimal way to scan
 - Table locks don't allow much concurrency but are cheaper to acquire and manage
- SQL Server locking strategy is dynamic
 - SQL Server chooses the lowest cost locking strategy (row, page, or table) at run time based upon input from the query optimizer



Lock Escalation

- Happens in two conditions:
 - 1) $>5,000$ locks
 - 2) lock memory $> 40 \%$
- Used to lower the number of locks taken by a transaction
 - Lock manager attempts to replace one transaction's many row or page locks with a single table-level lock
 - Escalation never converts row locks to page locks
 - There is no lock escalation on temporary tables or system tables.
- Lock “de-escalation” never occurs

Lock Time Out

- Application lock timeout
 - User configurable lock time-out
 - SET LOCK_TIMEOUT 10000
 - Rollback transaction
- Internal lock timeout
 - Internal lock request can time out
 - Does not rollback transaction

Locking Hints

- Can be specified using the SELECT, INSERT, UPDATE, and DELETE statements
- Direct SQL Server to the type of locks to be used
 - **Granularity hints:** ROWLOCK, PAGLOCK, TABLOCK
 - **Isolation Level hints:** HOLDLOCK, NOLOCK
 - READCOMMITTED, REPEATABLE, SERIALIZABLE, READUNCOMMITTED
 - **UPDLOCK:** use update lock rather than shared lock when reading
 - **XLOCK:** use exclusive lock instead
 - **READPAST:** will “skip” rows that are currently locked
- Used when a finer control of the types of locks acquired on an object is required
- Override the current transaction isolation level for the session

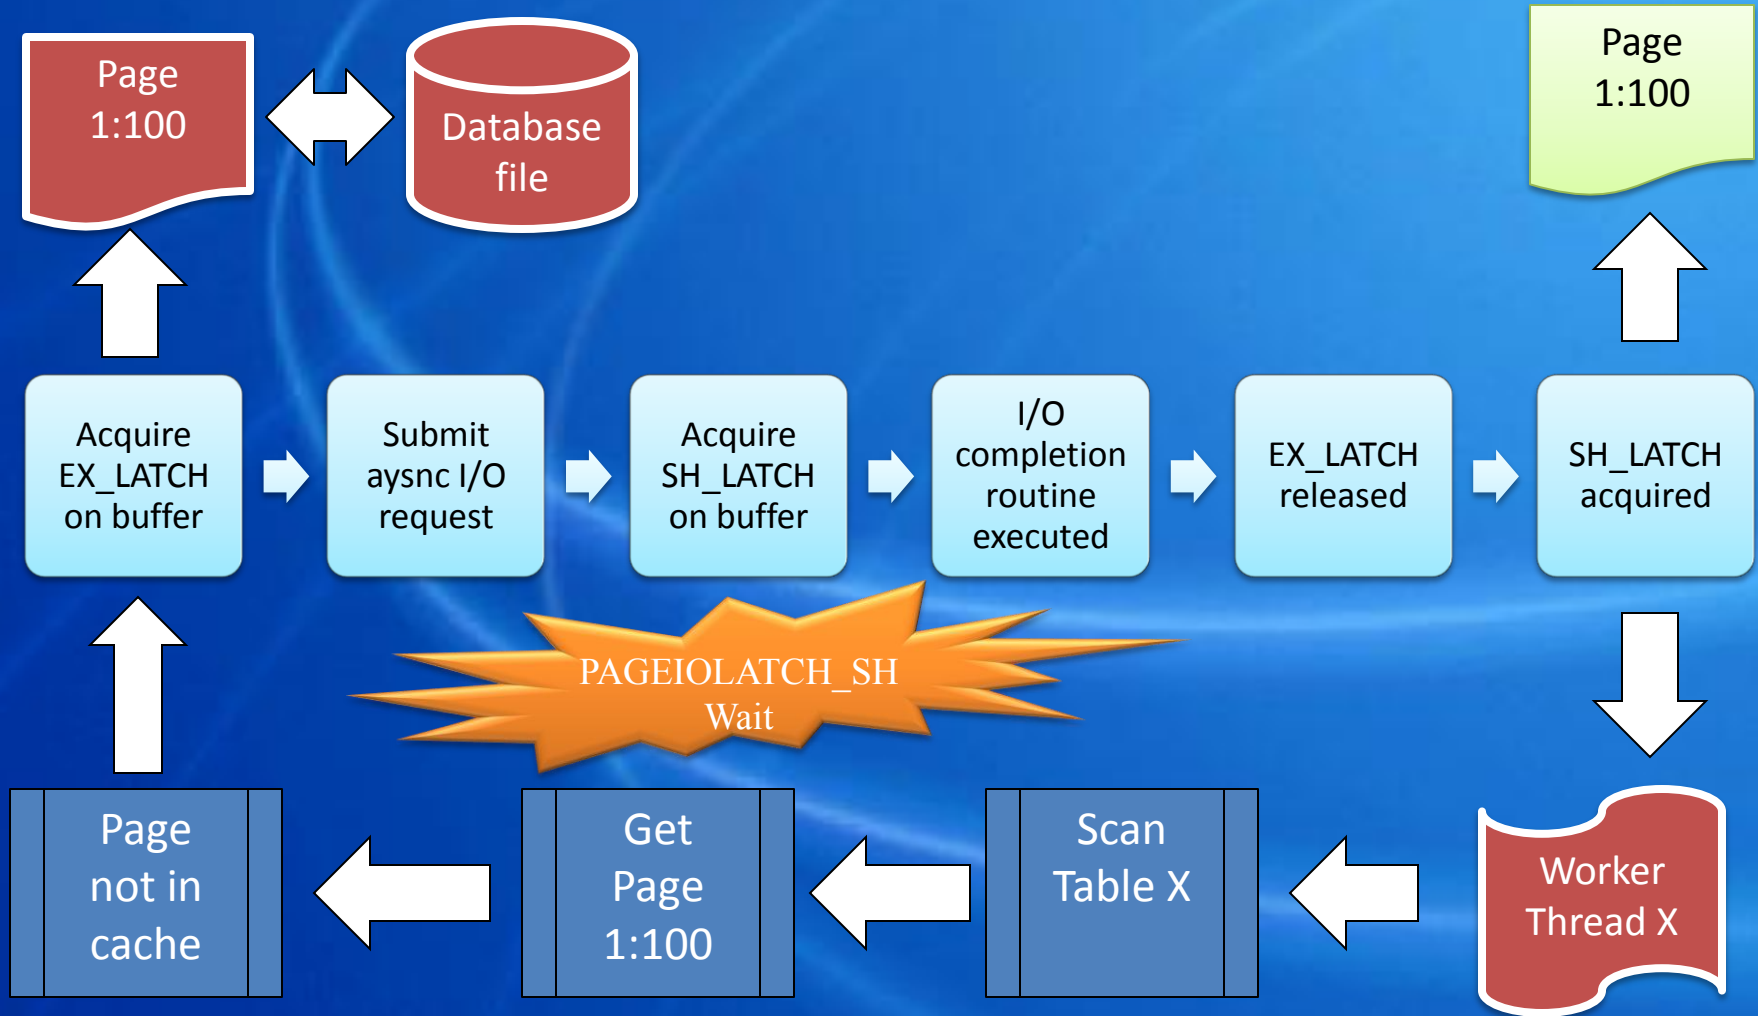
Recommendations to prevent locking issue

- Keep the transaction short and in one batch
- Separate readers from writers
- Beware of implicit transactions
- Use lowest Isolation Level required
- Process all Result Quickly
- Avoid Escalation when it is NOT necessary

PAGEIOLATCH_x

- Latches are short term synchronization objects. used to synchronize access to buffer pages. PageIOLatch is used for disk to memory transfers.
- If this is significant in percentage, it typically suggests memory pressure and disk IO subsystem issues. Check memory and disk counters.

How PAGEIOLATCH_x works for Reading a page from disk



PAGELATCH_x

- Occurs when a task is waiting for a latch for a buffer that is not in an I/O request.
- If this is significant in percentage, it typically indicates cache contention, not IO or memory.
- Gather blocking script. Try to find out the work load which has caused contention.
- Could consider gathering a dump file to double confirm your suspect.
- Review the known SQL issues.
- Also could consider using partitioned table

Who Needs a PAGELATCH_x When I Have Locks

```
INSERT VALUES
(3,300)
```

IX Page
100

EX_LATCH

```
INSERT VALUES
(4,400)
```

IX Page
100


EX_LATCH

Page 100

m_freedata=126

1	100	2	200
---	-----	---	-----

111 96

Page 100			
m_freedata=		156	
			
1	4		
141	126	111	96

Tempdb Contention

Where will I see latch contention?

- SGAM , PFS, and GAM pages
- System catalog pages

Why is there latch contention?

- Alloc pages must be latched as part of “transaction” of allocation
- System tables pages latched due to drop/create of a table

How do I reduce the contention?

- Create multiple data files
- All files have the same size
- Avoid autogrow and reduce tempdb usage

This could happen in a user database

- If users dropped/created tables frequently.

Demo

Problem:

The end user complains application “slow”.

Questions:

What do you think the problem is?

Why contention on tempdb occurs?

Action plan?

ASYNC_NETWORK_IO

- SQL Server has finished query execution within database engine.
- SQL Network I/O is asynchronized
- Normally caused by
 - network issue
 - Client does not fetch the data quickly
- Netmon trace is useful to identify the problem

Lab 1 - Blocking

- Learn how to analyze the blocking issue with blocking script and SQL trace.

Advanced blocking issues

- TokenAndPermUserStore
- Resource semaphore
- Resource semaphore query compile

TokenAndPermUserStore

Symptoms:

- Queries that typically run faster take a longer time to finish running
- CPU utilization for the SQL Server process is more than usual.
- When you experience decreased performance when you run an ad hoc query, you view the query from the **sys.dm_exec_requests** or **sys.dm_os_waiting_tasks** DMVs. However, the query does not appear to be waiting for any resource.
- The size of the TokenAndPermUserStore cache store is in the order of several hundred megabytes (MB)

```
SELECT SUM(single_pages_kb + multi_pages_kb) AS  
"CurrentSizeOfTokenCache(kb)" FROM sys.dm_os_memory_clerks  
WHERE name = 'TokenAndPermUserStore'
```

TokenAndPermUserStore

Solutions

- Explicitly parameterize ad hoc queries.
- Wrap ad hoc queries within stored procedures, and use stored procedures instead of directly executing ad hoc queries.
- Add the login that executes varied ad hoc queries as a member of the sysadmin server group.
- Flush entries from the TokenAndPermUserStore cache. -
DBCC FREESYSTEMCACHE ('TokenAndPermUserStore')

TokenAndPermUserStore

Solutions (cont.)

- SQL 2005 SP2: Trace Flag 4618 and 4610
- SQL 2005 SP3: Trace flag 4621 + add registry key

HKEY_LOCAL_MACHINE\SOFTWARE\Microsoft\Microsoft SQL Server\<MSSQL.X>\MSSQLServer

Name: TokenPermQuota

Type: DWORD

Value: <QuotaSize in Hex>

- SQL 2008: sp_configure - "access check bucket count" and "access check cache quota" option

Resource semaphore

- Memory grants were required for queries that performed sort or hash operations. These queries can be potentially blocked on RESOURCE_SEMAPHORE wait type—until the necessary grant could be acquired.
- If time-out, error 8645 is encountered.
- Acquiring the grant before execution ensures that the necessary memory will be available at runtime.

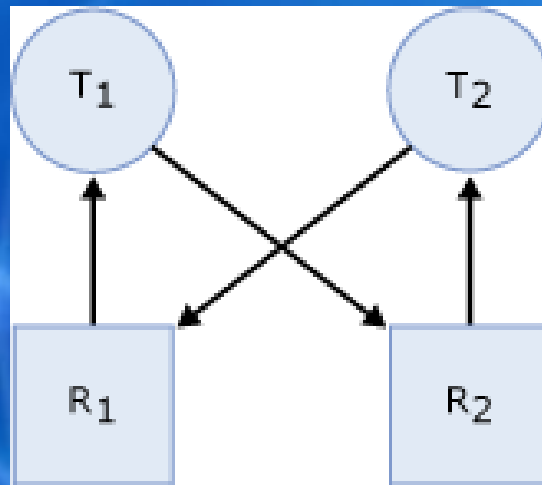
Resource semaphore query compile

- A large number of waits of type `RESOURCE_SEMAPHORE_QUERY_COMPILE` indicates a large number of concurrent compiles.
- SQL server has gateways for throttling individual query optimizations based on memory usage
- To prevent inefficient use of server resource.

Deadlock

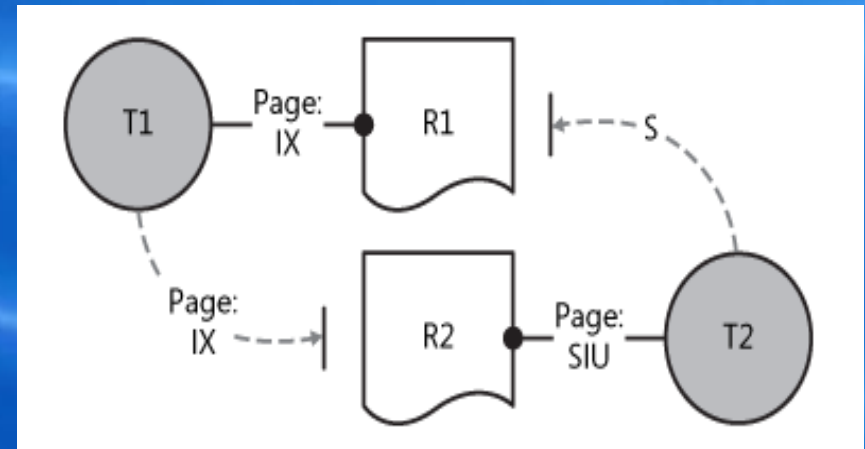
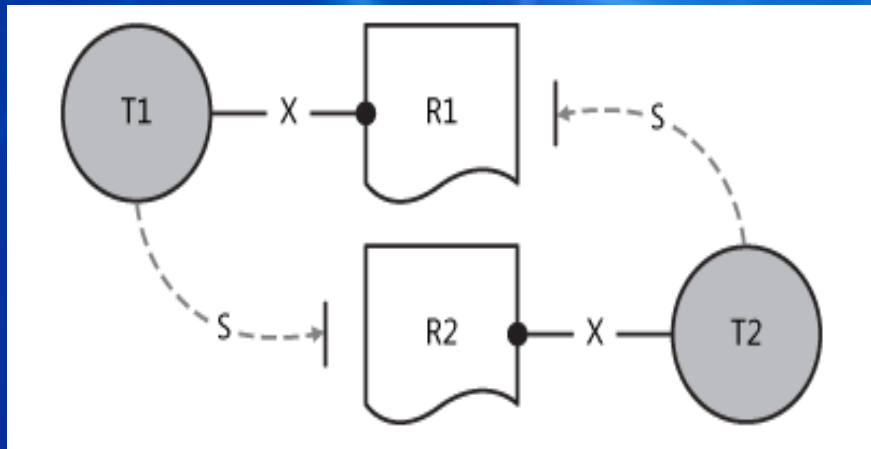
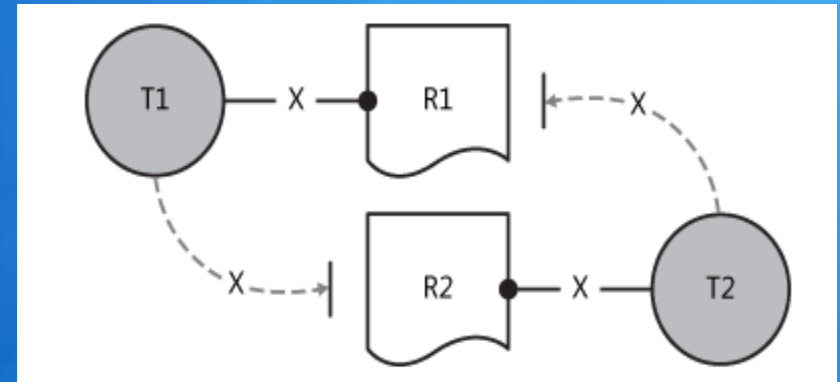
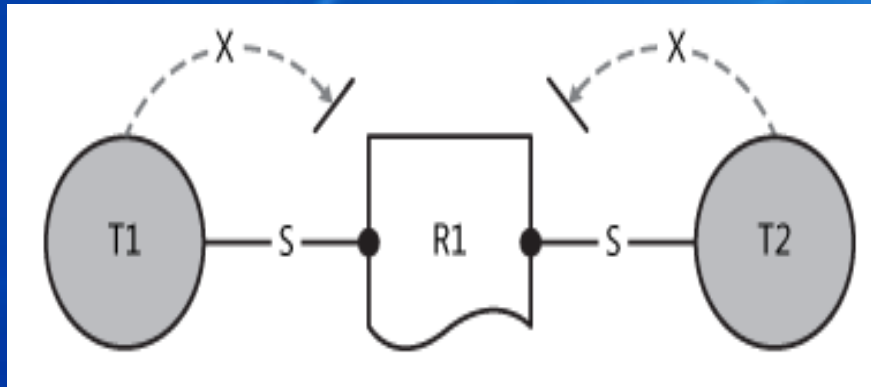
Deadlock

- Occurs when two connections are simultaneously waiting on a resource that the other connection holds



- Tran #1 holds Resource1 and is waiting on Resource2
- Tran #2 holds Resource 2 and is waiting on Resource1
- Each transaction is waiting on the other, and neither can proceed

Deadlock Examples



Deadlock Detection and resolution

- Lock manager detects deadlocks automatically by means of a background process called LOCK_MONITOR
- If deadlock is detected. Select the transaction that is cheapest to rollback - the deadlock victim
- Rollback the victim's transaction
- Notify the victim by means of a 1205 error:
Error 1205: Your transaction (process ID #%d) was deadlocked with another process and has been chosen as the deadlock victim. Rerun your transaction.

Information collection for deadlock

- Enable trace flag 1204 or 1222 (new from SQL Server 2005)
- SQL errorlog
- Capture SQL trace

How to avoid the deadlock

- Add the clustered index
- Reduce the lock holding time in the transaction
- Use the small transaction instead of the long duration transaction
- Commit the implicit transaction
- Reduce the lock range
- Change the lock type

Deadlock Analysis – Example Using Trace Flag 1204

Deadlock encountered Printing deadlock information

Wait-for graph

Node:1

KEY: 5:2121058592:1 (a70064fb1eac) CleanCnt:1 Mode: X Flags: 0x0

Grant List::

Owner:0x19165c00 Mode: X Flg:0x0 Ref:0 Life:02000000 SPID:58 ECID:0

SPID: 58 ECID: 0 Statement Type: UPDATE Line #: 1

Input Buf: Language Event: update authors set au_id = au_id

Requested By:

ResType:LockOwner Stype:'OR' Mode: U SPID:57 ECID:0 Ec:(0x1afcd520) Value:0x19167bc0 Cost:(0/10AC)

Node:2

KEY: 5:1977058079:1 (02014f0bec4e) CleanCnt:1 Mode: X Flags: 0x0

Grant List::

Owner:0x1916c220 Mode: X Flg:0x0 Ref:0 Life:02000000 SPID:57 ECID:0

SPID: 57 ECID: 0 Statement Type: UPDATE Line #: 1

Input Buf: Language Event: update titles set type = type

Requested By:

ResType:LockOwner Stype:'OR' Mode: U SPID:58 ECID:0 Ec:(0x1936d520) Value:0x1916bb80 Cost:(0/54)

Victim Resource Owner:

ResType:LockOwner Stype:'OR' Mode: U **SPID:58** ECID:0 Ec:(0x1936d520) Value:0x1916bb80 Cost:(0/54)

Lock Resource

Lock Mode

Useful if Stored Proc is Involved

Lock Mode

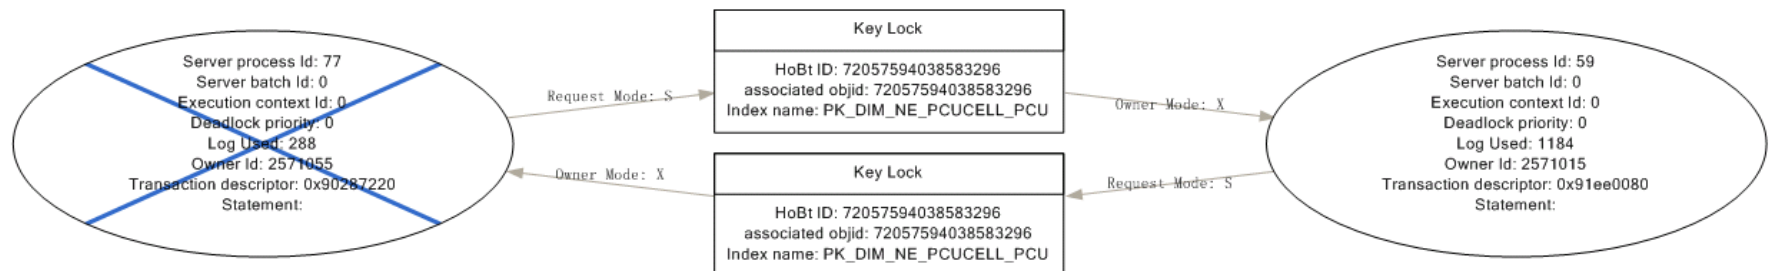
Deadlock Victim

Analyze deadlock in SQL 2005

- New trace flag is introduced: T1222
- New SQL trace event: Deadlock graph

SQL Server Profiler - [F:\...\deadlock\HP-DL580_WEDNESDAY_sp_trace.trc]

EventClass	BinaryData	DatabaseID	TransactionID	NTUserNa...	NTDomainN...	HostNa...	ClientProcessID	ApplicationName	LoginName	SPID	Duration	StartTime	EndTime
SP:StmtCompleted		10	2580661			HP-DL580	24232	DataImport	sa	65	3	2008-01-09 11:18:22.310	2008-01-09 11:18:22.313
SP:StmtCompleted		10	2580661			HP-DL580	24232	DataImport	sa	65	20	2008-01-09 11:18:22.293	2008-01-09 11:18:22.313
SQLTransaction		10	2580661			HP-DL580	24232	DataImport	sa	65	20	2008-01-09 11:18:22.293	2008-01-09 11:18:22.313
User Error Message		10				HP-DL580	24232	DataImport	sa	65		2008-01-09 11:18:22.327	
SQLTransaction		12	2580659			HP-DL580	24232	DataImport	sa	65	21	2008-01-09 11:18:22.293	2008-01-09 11:18:22.313
SP:StmtStarting		12	2580652			HP-DL580	24232	DataImport	sa	65		2008-01-09 11:18:22.390	
Lock:Deadlock Chain	0X0A000...	8	2571175						sa	4		2008-01-09 11:18:23.187	
Lock:Deadlock Chain	0X0A000...	8	2571015						sa	4		2008-01-09 11:18:23.203	
Lock:Deadlock Chain	0X0A000...	8	2571055						sa	4		2008-01-09 11:18:23.203	
SQLTransaction		1	2582872						sa	16	0	2008-01-09 11:18:23.310	2008-01-09 11:18:23.310
SQLTransaction		1	2582872						sa	16	0	2008-01-09 11:18:23.310	2008-01-09 11:18:23.310
SQLTransaction		1	2582876						sa	16	0	2008-01-09 11:18:23.327	2008-01-09 11:18:23.327
SQLTransaction		1	2582876						sa	16	0	2008-01-09 11:18:23.327	2008-01-09 11:18:23.327
Deadlock graph		1							sa	16		2008-01-09 11:18:23.327	
Showplan Statistics Profile	0X1A000...	12	2571175			HP-DL580	24232	DataImport	sa	65		2008-01-09 11:18:22.522	



Lab 2 - Deadlock

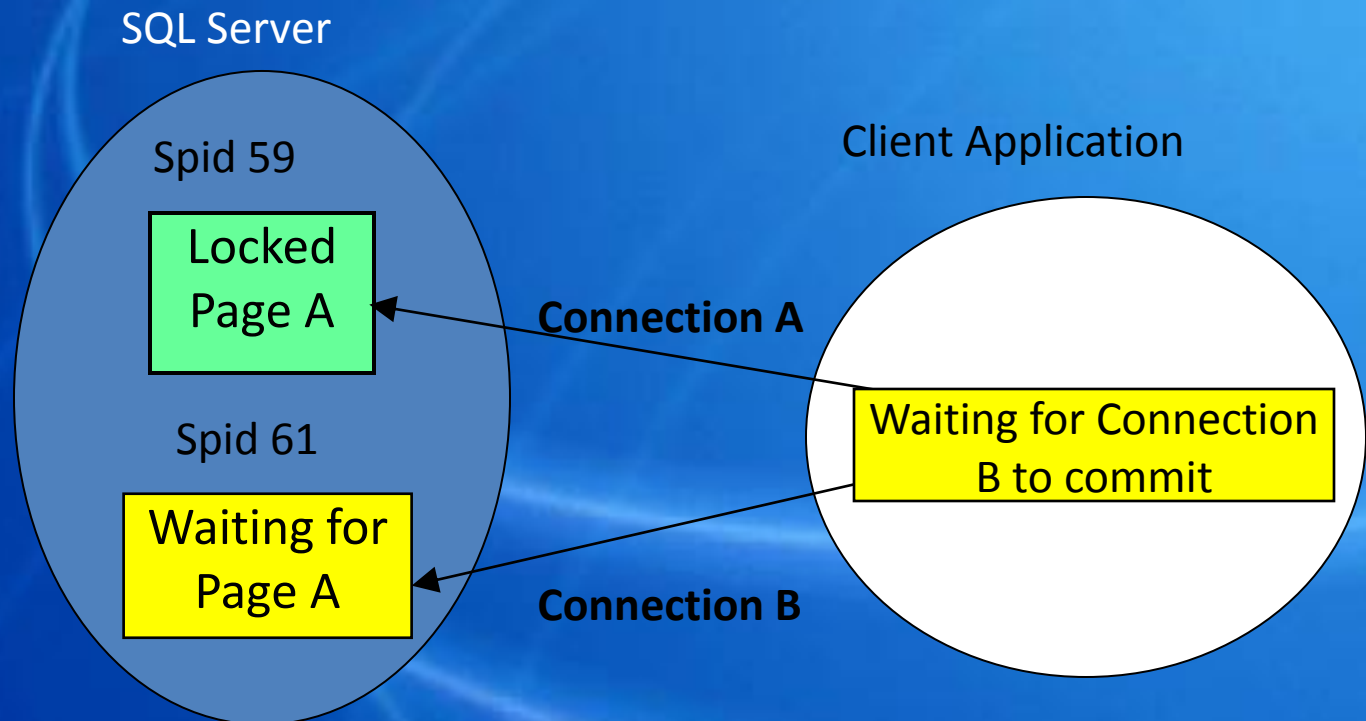
- Generate a deadlock issue manually
- Analyze the deadlock
- How to resolve it?

Distributed Deadlock

- Distributed deadlocks usually occur when one of the locked resources resides outside of SQL Server (or on another SQL Server)

Distributed Deadlock – Scenario 1

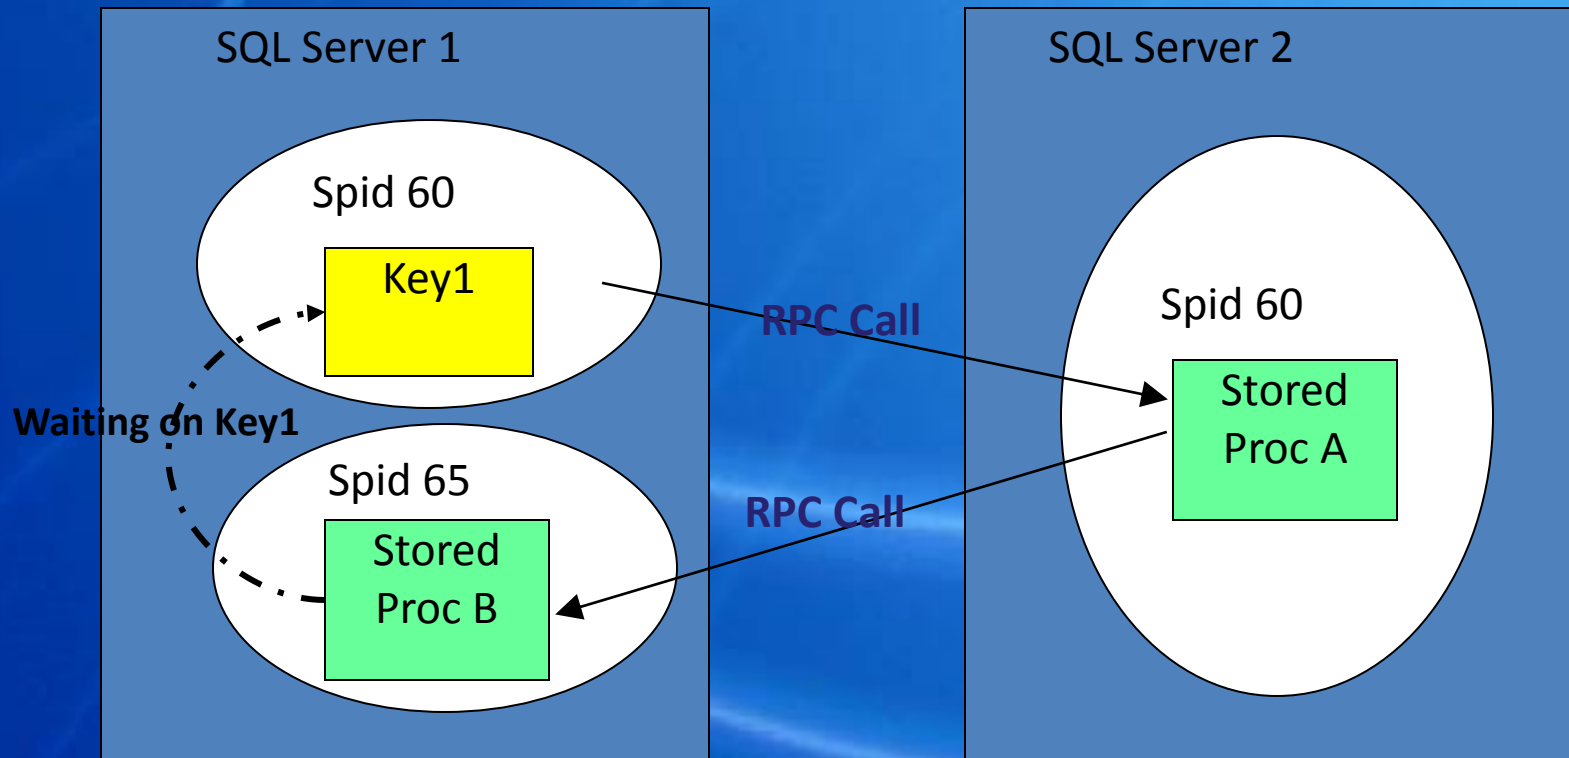
- Scenario 1



- This deadlock cannot be detected by SQL Server

Distributed Deadlock – Scenario 2

- Scenario 2



- This deadlock cannot be detected by SQL Server

Symptoms

- SQL server one or more SPIDs were waiting on ASYNC_NETWORK_IO
- Application server hang or ran slowly

Information to collect

- Netmon traces
- PSSDIAG/Blocking script
- Dumps for <application>.exe and sqlservr.exe

Review

- Blocking Issue
 - Locking
 - PageIOlatch
 - Pagelatch
 - ASYNC NETWORK IO
 - Tempdb Contention
- Deadlock

Q & A

Thank You!



"MICROSOFT CONFIDENTIAL. Distribution Only to Partners Under Nondisclosure. Microsoft makes no warranties, express or implied. ©2009 Microsoft Corporation. "

© 2008 Microsoft Corporation. All rights reserved. Microsoft, Windows, Windows Vista and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries. The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.